

## PCA analysis


- Standardization of the data

Variables measured at different scales do not contribute equally to the model fitting & model learned function and might create a bias. Thus, to deal with this potential problem feature-wise standardized ( $\mu=0$ ,  $\sigma=1$ ) is usually used before model fitting.

```
from sklearn.preprocessing import StandardScaler
#object
scaler = StandardScaler()
# Fit on training set only.
scaler.fit(X_train)
# Apply transform to both the training set and the test set.
X_train = scaler.transform(X_train)
X_test = scaler.transform(X_test)
X = scaler.transform(X)
```


- Principal Component Analysis

Without setting the value of `n_components`:

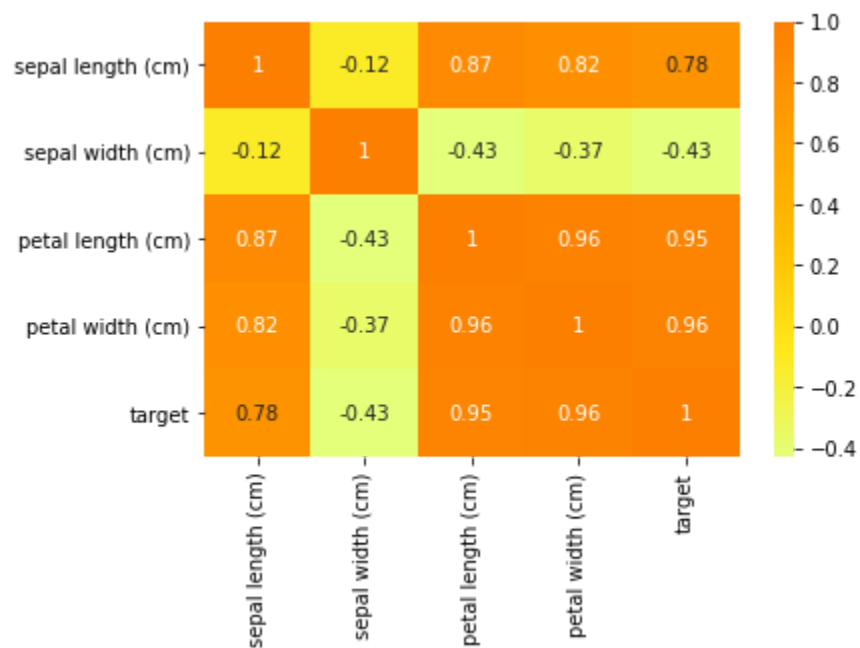


The variance captured by the first 2 features when arranged in descending order contribute to more than 90% of the total variance by the features. Thus, we can eliminate the remaining 2 features as they don't contribute much to the overall variance.


- Dimensions that are primary contributors to the first eigenvector



The above correlation plot shows that the primary contributors to the first eigenvector petal length (cm) and sepal length (cm), both are positively correlated to the eigenvector and each other as follows:





- The plot of the transformed data using the first two eigenvectors is as follows:

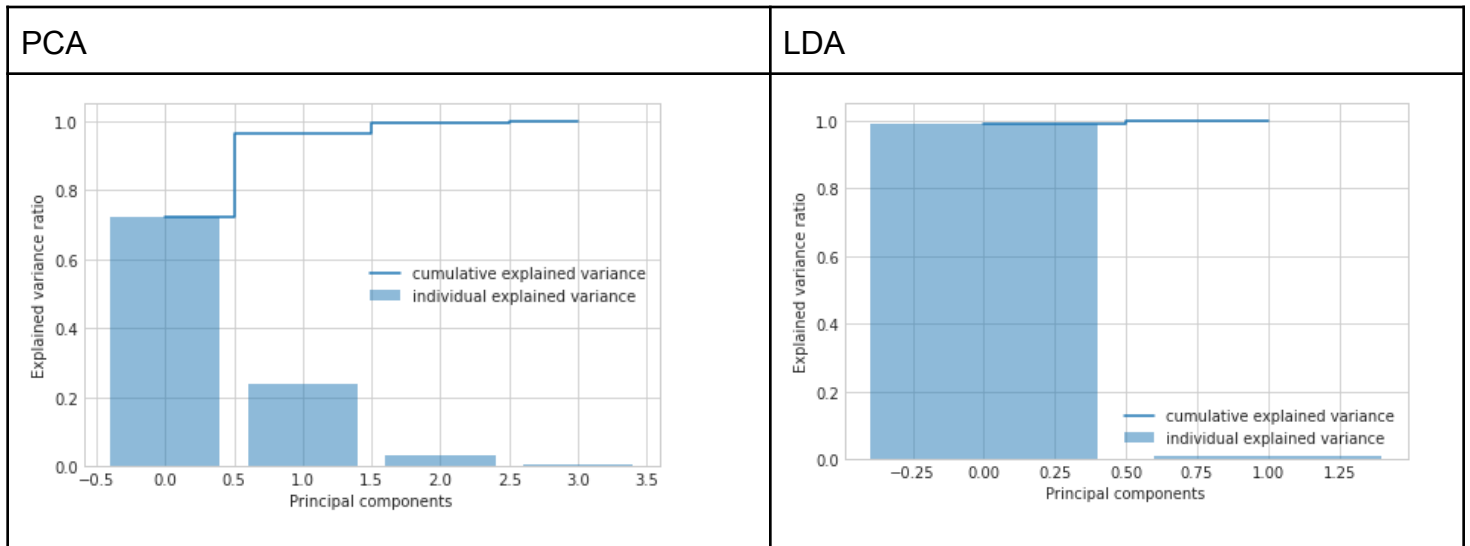


#### PCA and LDA comparison

- Plots for distribution of samples using the first 2 principal components and the first 2 linear discriminants:



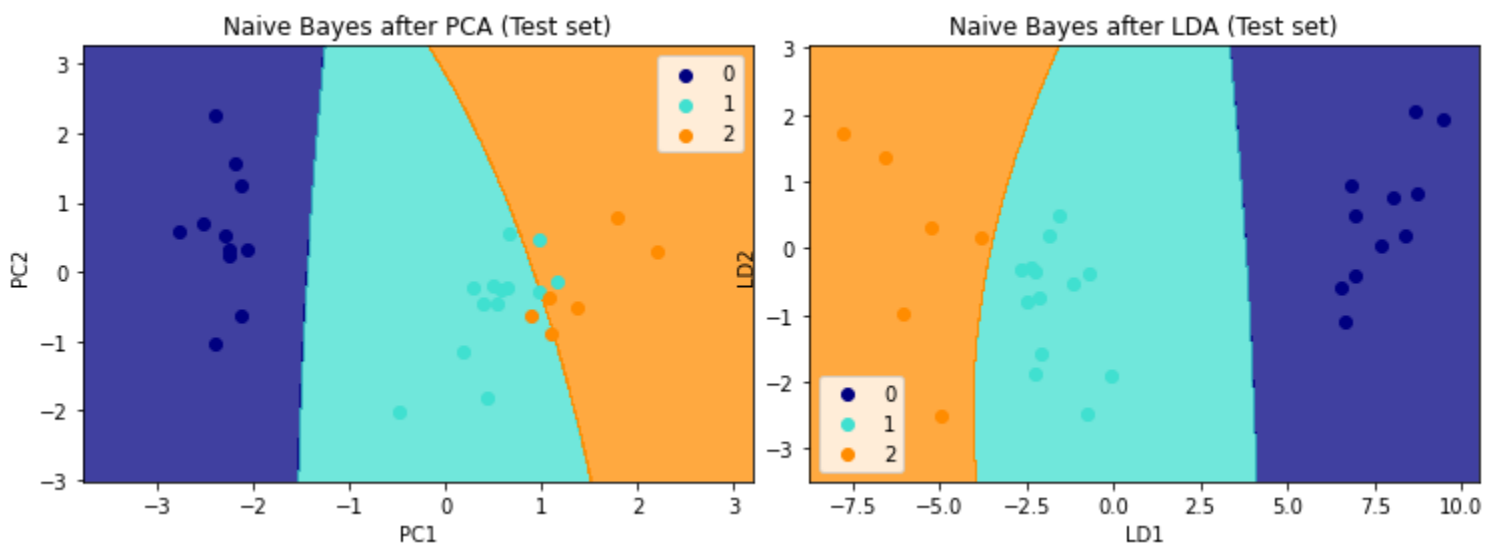
- Comparing PCA and LDA



- Learn a Bayes classifier using the original features and compare its performance with the features obtained in part:

|                    | Without Dimensionality reduction   | PCA  | LDA  |
|--------------------|--|--|--|
| score()            | 0.9666666666666667   | 0.8666666666666667   | 1.0  |
| Confusion Matrices | <pre>array([[11,  0,  0],        [ 0, 13,  0],        [ 0,  1,  5]])</pre> | <pre>array([[11,  0,  0],        [ 0, 11,  2],        [ 0,  2,  4]])</pre> | <pre>array([[11,  0,  0],        [ 0, 13,  0],        [ 0,  0,  6]])</pre> |

- Visualizing the test results:



## Feature Selection

- Preprocessing and Exploratory Data Analysis:



```
df.shape

(768, 9)

df.columns

Index(['preg', 'plas', 'pres', 'skin', 'test', 'mass', 'pedi', 'age', 'class'], dtype='object')

df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 768 entries, 0 to 767
Data columns (total 9 columns):
#   Column  Non-Null Count  Dtype
---  -
0    preg    768 non-null    int64
1    plas    768 non-null    int64
2    pres    768 non-null    int64
3    skin    768 non-null    int64
4    test    768 non-null    int64
5    mass    768 non-null    float64
6    pedi    768 non-null    float64
7    age     768 non-null    int64
8    class   768 non-null    int64
dtypes: float64(2), int64(7)
memory usage: 54.1 KB
```


df.describe()

|       | preg       | plas       | pres       | skin       | test       | mass       | pedi       | age        | class      |
|-------|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean  | 3.845052   | 120.894531 | 69.105469  | 20.536458  | 79.799479  | 31.992578  | 0.471876   | 33.240885  | 0.348958   |
| std   | 3.369578   | 31.972618  | 19.355807  | 15.952218  | 115.244002 | 7.884160   | 0.331329   | 11.760232  | 0.476951   |
| min   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.000000   | 0.078000   | 21.000000  | 0.000000   |
| 25%   | 1.000000   | 99.000000  | 62.000000  | 0.000000   | 0.000000   | 27.300000  | 0.243750   | 24.000000  | 0.000000   |
| 50%   | 3.000000   | 117.000000 | 72.000000  | 23.000000  | 30.500000  | 32.000000  | 0.372500   | 29.000000  | 0.000000   |
| 75%   | 6.000000   | 140.250000 | 80.000000  | 32.000000  | 127.250000 | 36.600000  | 0.626250   | 41.000000  | 1.000000   |
| max   | 17.000000  | 199.000000 | 122.000000 | 99.000000  | 846.000000 | 67.100000  | 2.420000   | 81.000000  | 1.000000   |

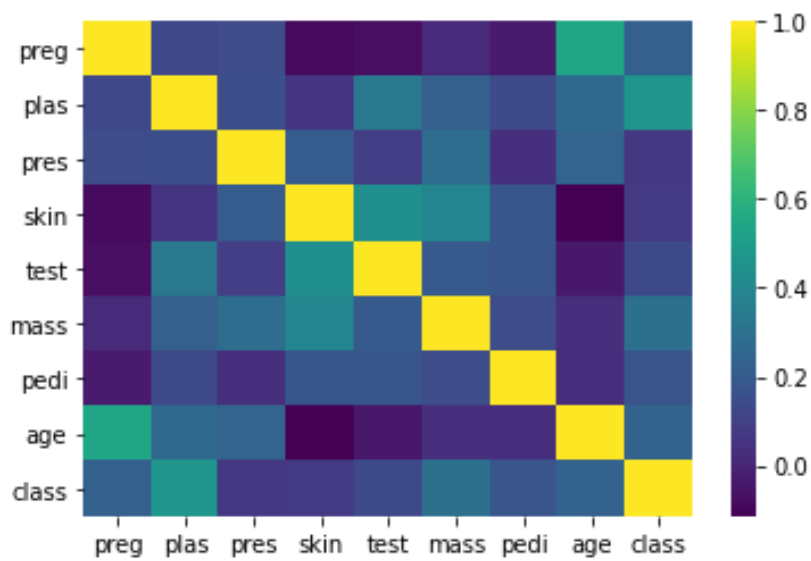
```
df.isnull().sum()

preg      0
plas      0
pres      0
skin      0
test      0
mass      0
pedi      0
age       0
class     0
dtype: int64
```

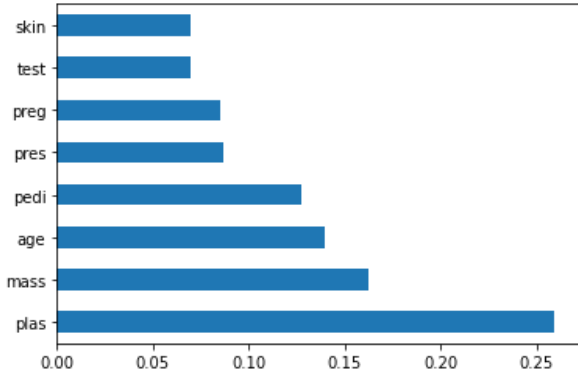
The plot for missing data:



Correlation Plot:



- Features having high significance using both of the methods

| Model                        | Feature Scores   | Highly significant              |
|------------------------------|--|---------------------------------|
| Chi-squared statistical test | <pre> Features      Score test  2175.565273 plas  1411.887041 age   181.303689 mass  127.669343 preg  111.519691 skin   53.108040 pres   17.605373 pedi    5.392682           </pre> | ['test', 'plas', 'age', 'mass'] |
| Random Generation            |    | ['plas', 'mass', 'pedi', 'age'] |


- Accuracies and F1 Scores based comparison

It turns out that the model trained with all features performed better than the other two while the models with selected features perform with similar efficiencies.

| Model                        | F1 Score           | Accuracy           |   |
|------------------------------|--------------------|--------------------|---|
| With all features            | 1.0                | 1.0                | <pre>array([[107,  0],        [  0, 47]])</pre> |
| Chi-squared statistical test | 0.6363636363636364 | 0.7922077922077922 | <pre>array([[97, 10],        [22, 25]])</pre>   |
| Random Generation            | 0.6097560975609756 | 0.7922077922077922 | <pre>array([[94, 13],        [19, 28]])</pre>   |




- Correlation Matrix



- Compute correlated features with a threshold of 70%.

We have to find out the correlation between the features and remove the features which have a correlation coefficient greater than a certain limit (here 70% is the threshold)

Plot after applying the threshold:



- This shows that all the features are correlated with a threshold of 70%

\_\_\_\_\_

```
features_c = corr.columns
for p in corr.values:
    for i in range(len(features_c)):
        if(p[i]< 0.7 and p[i]>= -0.7):
            p[i] = 0
corr
```

[illegible]